

CIS 5511. Programming Techniques

Graphs (2)

1. Minimum spanning trees

For an undirected and connected graph $G = (V, E)$, its *spanning tree* is a subgraph $G' = (V, E')$, which is still connected and $|E'| = |V| - 1$, so G' does not have any cycle. This type of *tree* is different from the ones introduced previously, as there is no root specified.

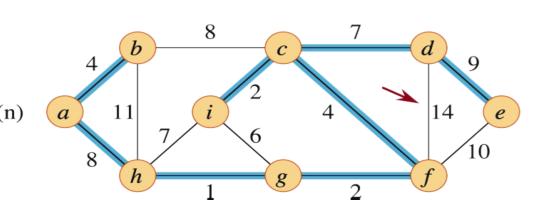
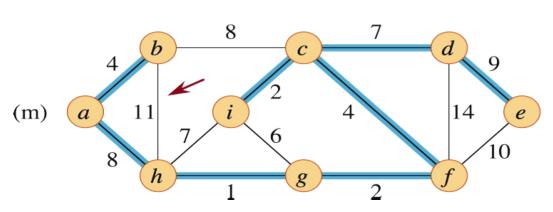
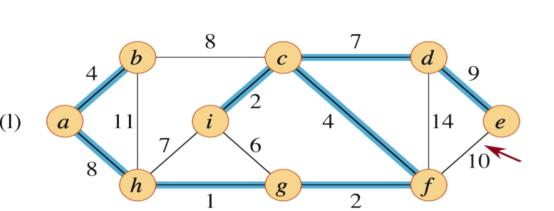
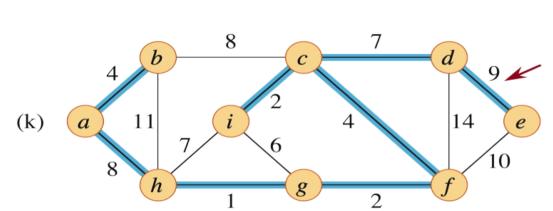
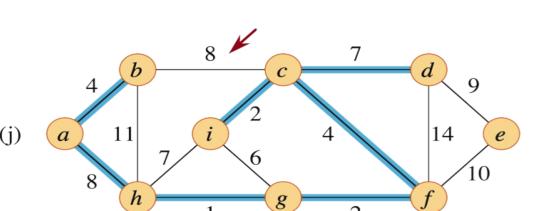
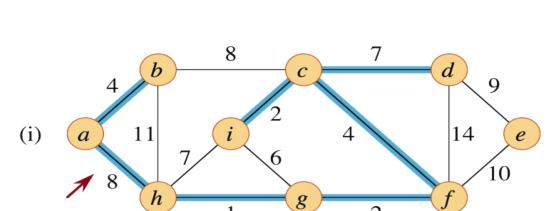
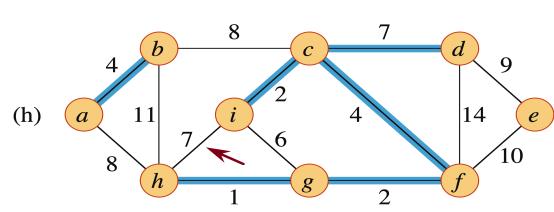
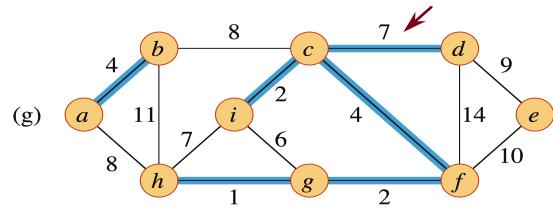
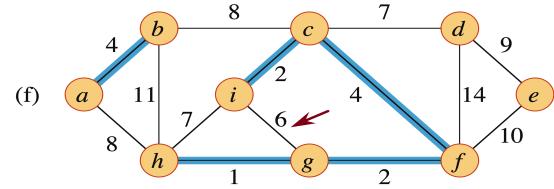
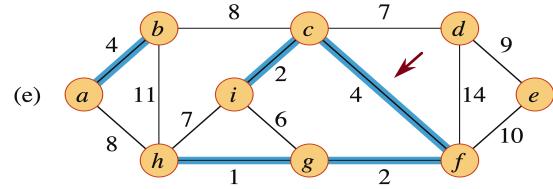
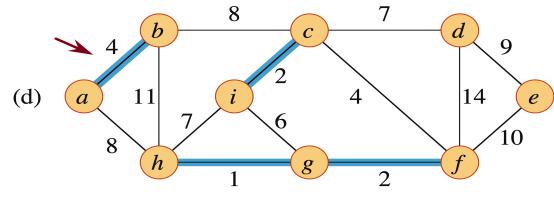
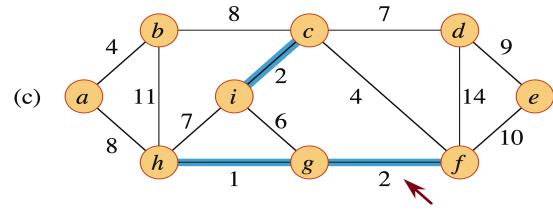
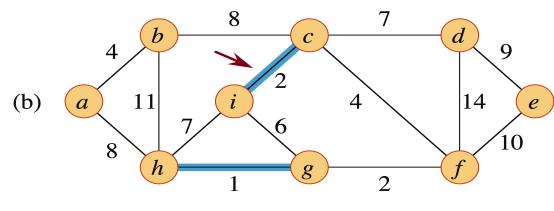
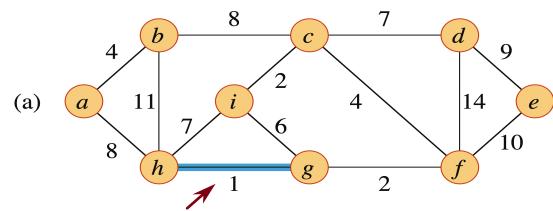
If G is weighted, its *minimum spanning tree* has a minimum value of total weight $\sum w(u, v)$ [for $(u, v) \in E'$] among all spanning trees of G .

Kruskal's algorithm each time adds an edge with the least w , and connects two previously unconnected subgraphs. The algorithm uses a set to represent the vertices in a subtree, and $\text{Find-Set}(u)$ identifies the set in which vertex u belongs.

MST-KRUSKAL(G, w)

```
1   $A = \emptyset$ 
2  for each vertex  $v \in G.V$ 
3      MAKE-SET( $v$ )
4  create a single list of the edges in  $G.E$ 
5  sort the list of edges into monotonically increasing order by weight  $w$ 
6  for each edge  $(u, v)$  taken from the sorted list in order
7      if  $\text{FIND-SET}(u) \neq \text{FIND-SET}(v)$ 
8           $A = A \cup \{(u, v)\}$ 
9          UNION( $u, v$ )
10 return  $A$ 
```

Example:



Kruskal's algorithm is a greedy algorithm, and takes $O(E \lg E)$ time, which is the same as $O(E \lg V)$.

Prim's algorithm adds vertices into the tree one by one, according to their distance to the tree built so far, starting from a root r . Q is a priority queue of vertices, where for each vertex v , $\text{key}[v]$ is the minimum weight of any edge connecting v to a vertex in the tree.

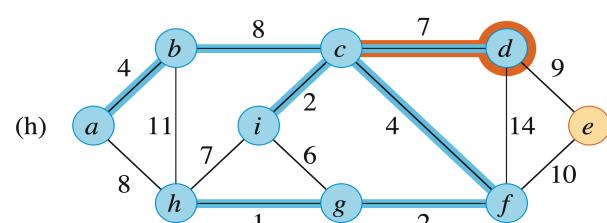
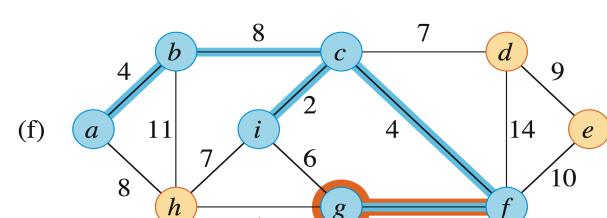
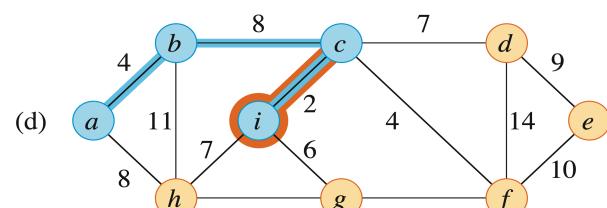
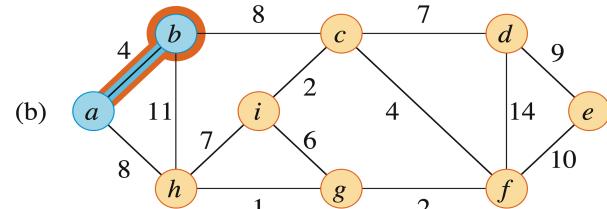
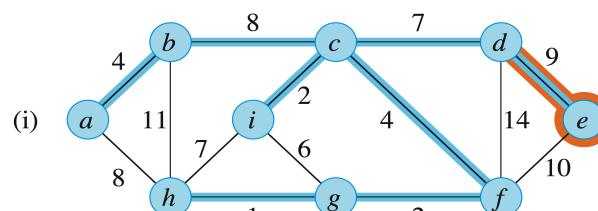
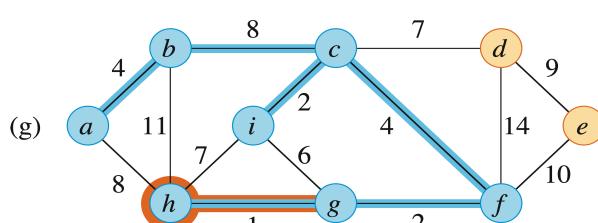
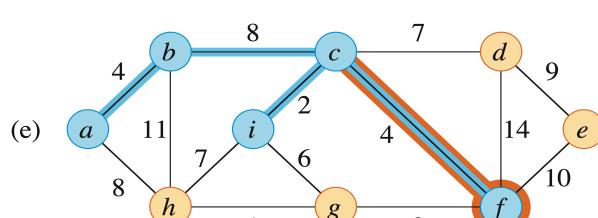
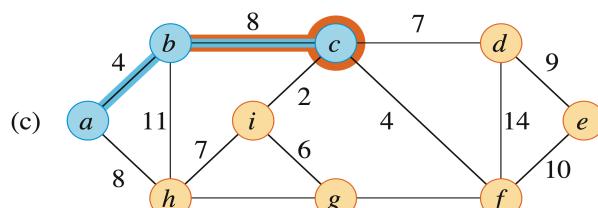
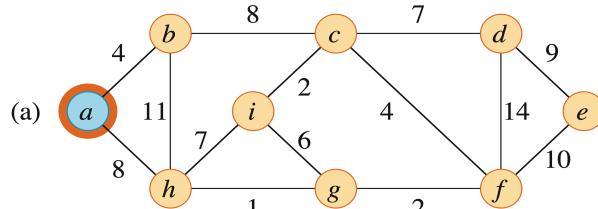
MST-PRIM(G, w, r)

```

1  for each vertex  $u \in G.V$ 
2     $u.key = \infty$ 
3     $u.\pi = \text{NIL}$ 
4   $r.key = 0$ 
5   $Q = \emptyset$ 
6  for each vertex  $u \in G.V$ 
7    INSERT( $Q, u$ )
8  while  $Q \neq \emptyset$ 
9     $u = \text{EXTRACT-MIN}(Q)$  // add  $u$  to the tree
10   for each vertex  $v$  in  $G.Adj[u]$  // update keys of  $u$ 's non-tree neighbors
11     if  $v \in Q$  and  $w(u, v) < v.key$ 
12        $v.\pi = u$ 
13        $v.key = w(u, v)$ 
14       DECREASE-KEY( $Q, v, w(u, v)$ )

```

Example:



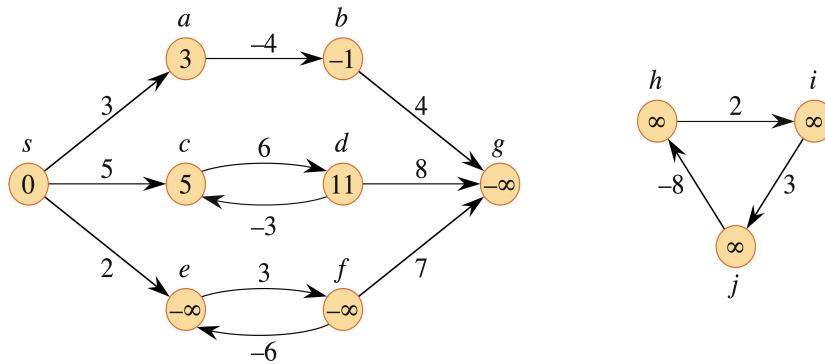
Prim's algorithm also takes $O(E \lg V)$ time.

2. Single-source shortest paths

In a weighted graph, one type of optimization problem is to find the shortest path between vertices, and it can be one-to-one, one-to-many, or many-to-many. In particular, "single-source" means to find the shortest paths from one source vertex to all the other vertices, so is the one-to-many case. Here "distance" or "weight" can represent many different measurements, so can have any finite (positive, negative, or zero) value.

The weight of a path is the sum of the weights of the edges in the path. A shortest path is a path whose weight has the lowest value among all paths with the same source and destination. A section of a shortest path is also a shortest path.

In general, a shortest path can be defined in graphs that are either directed or undirected, with or without cycle. However, a shortest path cannot contain cycle. Therefore, in graph $G = (V, E)$, the shortest path contains less than $|V|$ edges. For example, in the following graph the lengths of the shortest paths from s to each vertex is marked, which include positive and negative infinite as special cases.



A trivial solution to the problem is to traverse the graph and enumerate all possible paths, and compare their distances. However, this solution is usually too time-consuming to be useful.

If all edges are equally-weighted, the breadth-first search algorithm BFS introduced previously finds shortest paths from a source vertex to all other vertices. However, it does not work if the edges have different weights.

To represent the shortest paths, an array d is used to record the shortest distance found so far from the source to each vertex, and an array π for the predecessor of each vertex on that path. The following algorithm initializes the data structure.

```
INITIALIZE-SINGLE-SOURCE( $G, s$ )
1 for each vertex  $v \in G.V$ 
2    $v.d = \infty$ 
3    $v.\pi = \text{NIL}$ 
4  $s.d = 0$ 
```

Many shortest-path algorithms use the "relaxation" step, which maintains the shortest paths from a given source to v found so far, and try to improve them when a new vertex u is taken into consideration. The matrix w stores the weights of the edges.

$\text{RELAX}(u, v, w)$

```

1  if  $v.d > u.d + w(u, v)$ 
2       $v.d = u.d + w(u, v)$ 
3       $v.\pi = u$ 

```

There are various ways to use this step to get the shortest paths, though it is important to see that after the distance to u is undated, the same step may need to be repeated.

Bellman-Ford algorithm processes the graph $|V| - 1$ passes. In each pass, the edges are tried one-by-one to relax the distance. After that, if there is still a possible relaxation, the graph must contain negatively-weighted cycle, so there is no shortest path to the involved vertices.

$\text{BELLMAN-FORD}(G, w, s)$

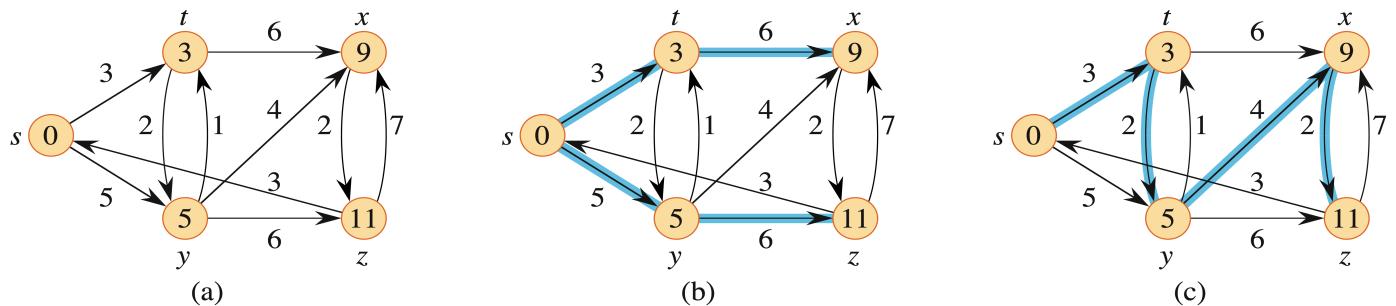
```

1  INITIALIZE-SINGLE-SOURCE( $G, s$ )
2  for  $i = 1$  to  $|G.V| - 1$ 
3      for each edge  $(u, v) \in G.E$ 
4          RELAX( $u, v, w$ )
5  for each edge  $(u, v) \in G.E$ 
6      if  $v.d > u.d + w(u, v)$ 
7          return FALSE
8  return TRUE

```

The running time is $O(V E)$.

Example:



If the graph is a DAG, there are faster solutions. The following algorithm topologically sorts the vertices first, then determine the shortest paths for each vertex in that order. It runs in $\Theta(V + E)$ time, which comes from topological sorting (and DFS).

Example:

Dijkstra's algorithm works for weighted graphs without any negative weight. It repeatedly selects the vertex with the shortest path, and uses it to relax the paths to other vertices.

This is a greedy algorithm. Its running time is $O((V + E) \lg V)$.

Example:

If only the shortest path to a single destination is needed, stop the algorithm at that vertex. The worst-case complexity will be the same.

3. All-pairs shortest paths

If we want the shortest paths between every pair of vertices, it is inefficient to repeat an algorithm for single-source with each vertex as source.

Many algorithms extend the weight matrix W to L that records the shortest paths found so far, plus a predecessor matrix Π , where π_{ij} is the predecessor of j on some shortest path from i to j , i.e., the last stop. Given it, the following algorithm prints the shortest path.

The following dynamic-programming algorithm extends shortest paths starting with single edges, and in each step tries to add one more edge. Here $n = |V|$.

Each call to the above algorithm will extend the length of path under consideration by one edge. Therefore, if we start with $L = W$, and repeatedly call the algorithm $n - 1$ times, we will get a matrix for the shortest paths. This algorithm takes $\Theta(n^4)$ time:

Example:

The above algorithm can be improved by updating the loop variable not as $m = m + 1$ but as $m = 2m$, and change the W in line 5 to $L^{(m-1)}$, so as to achieve $\Theta(n^3 \lg n)$ time.

Floyd-Warshall algorithm solves the problem by adding one possible intermediate vertex into the shortest paths in each step:

For the same problem, the intermediate results are:

The running time of the above algorithm is $\Theta(n^3)$, so is more efficient than the above two.

A similar algorithm calculates the transitive closure of a graph, where $T^{(n)}_{ij} = 1$ if and only if there is a path from i to j .